

# Project 1 — Tiny GPT-2 Q&A; Chatbot

## Why This Project Matters

This is your “Hello World” of fine-tuning. GPT-2 is small enough to run on Google Colab (free tier) or a low-end GPU. You’ll learn the entire workflow: dataset prep → model loading → training → inference. Skills from here carry forward to bigger models.

## Prerequisites

- Basic Python (loops, functions, variables)
- Hugging Face Transformers installation: `pip install transformers datasets accelerate`
- What is a tokenizer? Tokenizers split text into tokens the model understands.
- GPU access (Colab: Runtime → Change runtime type → GPU)

## Dataset

For Q&A, we’ll use the SQuAD v2 dataset from Hugging Face.

Link: [https://huggingface.co/datasets/squad\\_v2](https://huggingface.co/datasets/squad_v2)

Size: ~150k question-answer pairs.

Download & load:

```
from datasets import load_dataset
dataset = load_dataset("squad_v2")
print(dataset)
```

## Step-by-Step Build

Step 1: Load tokenizer & model:

```
from transformers import GPT2Tokenizer, GPT2LMHeadModel
model_name = "gpt2"
tokenizer = GPT2Tokenizer.from_pretrained(model_name)
model = GPT2LMHeadModel.from_pretrained(model_name)
```

Step 2: Preprocess dataset:

```
def preprocess(example):
    prompt = f"Question: {example['question']} Answer: {example['answers']['text'][0]} if example
    tokens = tokenizer(prompt, truncation=True, padding="max_length", max_length=128)
    tokens["labels"] = tokens["input_ids"].copy()
    return tokens
```

```
tokenized_dataset = dataset.map(preprocess, batched=True)
```

Step 3: Training loop:

```
from transformers import Trainer, TrainingArguments
training_args = TrainingArguments(
    output_dir="./gpt2-qa",
    evaluation_strategy="epoch",
    learning_rate=5e-5,
    per_device_train_batch_size=2,
    num_train_epochs=3,
    weight_decay=0.01,
```

```

        push_to_hub=False
    )
    trainer = Trainer(
        model=model,
        args=training_args,
        train_dataset=tokenized_dataset["train"],
        eval_dataset=tokenized_dataset["validation"]
    )
    trainer.train()

```

Step 4: Save & test:

```

trainer.save_model("./gpt2-qa")
input_text = "Question: What is fine-tuning?"
inputs = tokenizer(input_text, return_tensors="pt")
outputs = model.generate(**inputs, max_length=50)
print(tokenizer.decode(outputs[0]))

```

## Design Thinking

- GPT-2 small fits in <2GB VRAM, fast iterations. - Concatenating Q&A; into one prompt trains GPT-2 on our format. - 128 tokens keeps training fast while covering most questions and answers.

## Troubleshooting

Problem | Cause | Fix

CUDA out of memory | Batch too big | Reduce per\_device\_train\_batch\_size

Tokenizer error | Missing pad token | tokenizer.pad\_token = tokenizer.eos\_token

Model outputs nonsense | Not enough epochs | Increase num\_train\_epochs

## Flowchart

[Dataset] → [Preprocess Q&A] → [Tokenization] → [Model Load] → [Training Loop] → [Save Model]

## Mini Quiz

1. Why set labels = input\_ids in fine-tuning?
2. Purpose of truncation?
3. Why start with GPT-2 small?

Answers:

1. Model learns to predict next tokens.
2. Prevents input from exceeding max length.
3. Faster and less VRAM heavy.

## Side Quest

- Use your own Q&A dataset.
- Try distilgpt2 for faster runs.
- Experiment with temperature in generate() for creative answers.